

InboxIQ

Email Intelligence for Startup Founders

Jinxuan Lyu | Agentic AI | Spring 2026

~\$0.001

per email cost

~85%

gross margin

\$29

price / month

~4 users

payback at

1. The User

Early-stage startup founders, pre-Series A, typically 1–10 employees, who are the primary point of contact for investors, customers, partners, and candidates. They receive 50–150 meaningful emails per week and personally handle all high-stakes correspondence.

Concretely: a Columbia MBA founder running a seed-stage B2B SaaS company. They may spend 2–3 hours per day on email. Every investor reply, customer complaint, or partnership inquiry carries real business consequence. They cannot afford to miss signals or respond poorly, but they also cannot afford to think slowly.

2. The Problem

Today, founders handle every email the same way: open, read, decide how urgent it is, google the sender, think through a response, and write. This process is slow, inconsistent, and context-poor.

Three specific pain points:

- Triage is manual. Founders scan subject lines and senders to assess priority, often getting it wrong or spending mental energy on low-value emails (newsletters, cold sales outreach).
- Context is scattered. Replying to an investor means googling them, checking LinkedIn, and recalling prior conversations before writing a single word.
- Drafting is repetitive. Most email situations (investor check-ins, customer complaints, recruiting outreach) follow recognizable patterns, yet founders start from scratch every time.

Existing tools (Gmail, Superhuman, Notion AI) offer spell-check and sorting. None understand the email well enough to tell you who sent it, what they actually want, how urgent it is relative to today's date, and give you three ready-to-send replies with different strategic angles. InboxIQ does all four in one step.

3. The Economics

Revenue model

SaaS subscription. \$29/month per founder (individual), \$79/month per team (up to 5). Target: 500 individual users at month 12.

Cost to serve — one user-month

Item	Assumption	Monthly Cost
Gemini 2.5 Flash (full pipeline)	100 emails/day × 30 days × \$0.0013	~\$3.90
Gemini 2.5 Flash (newsletter skip)	30% of emails skip research step	−\$0.36 saved

Cloud Run compute	~10s per request, minimal idle	~\$0.50
Google Search grounding API	~60 searches/day (non-newsletter)	~\$0.18
Total cost to serve		~\$4.22 / user-month
Price		\$29 / user-month
Gross margin		~85%

Where it breaks

- At <15 emails/day per user, margin improves further (cost drops linearly, price is fixed).
- At >300 emails/day, cost approaches \$12/month — still 59% margin, but warrants tiered pricing.
- Google Search grounding is the most volatile cost item. If disabled for all emails, cost drops to ~\$2.40/month and margin rises to 92% — a configurable fallback already built into the product.
- Break-even on infrastructure fixed costs (~\$50/month Cloud Run minimum) at 4 paying users.

4. Why These Technical Choices

Every technical decision was made to directly serve the user's problem and the economics.

Google ADK multi-agent	The user's emails span wildly different contexts — an investor email needs web research; a newsletter needs none. A single monolithic prompt cannot adapt. ADK's LLMagent orchestration lets each specialist agent (triage, research, drafter) do one thing well. The model decides the path at runtime — that's the product's intelligence.
Conditional routing (skip research)	Research costs money (~\$0.003/call with search). Newsletter and sales outreach emails don't benefit from sender research — the founder knows who The AIWeekly Team is. Skipping research for these categories cuts ~30% of API cost at no quality cost. This routing logic is the direct reason the unit economics work.
get_current_date FunctionTool	Founders forward old emails, reference past conversations, and ask 'is this still urgent?' A model with a fixed training cutoff cannot assess 'end of this week' without knowing today's date. The date tool + UI date picker let users set the reference date to when the email was actually sent — making urgency assessment accurate across time.
Gemini 2.5 Flash (not Pro)	Flash is 10x cheaper than Pro with comparable quality for structured extraction and reply drafting. At \$0.0013/email vs ~\$0.013/email for Pro, the margin difference is 85% vs 55%. The model choice is a direct business decision.
Google Search grounding	Founders reply better when they know who they're talking to. A cold search for 'Roelof Botha Sequoia' gives the model current context — recent investments, fund stage, public statements — that personalizes the reply in ways that matter. This is the product's key differentiator from a generic AI email tool.
Structured output (JSON schemas)	The pipeline has three agents feeding each other. If triage returns freeform text, the drafter cannot reliably extract 'category: investor' to decide tone. Strict JSON schemas enforced via system prompts make agent chaining deterministic. The product works consistently, not just in demos.